

Moving the Mean Toward the Known Good, Not Beyond It: What Inference-Time Interventions and Weight Consolidation Buy in Open-Ended Generation

Roberto I. Ono Filho
Independent researcher
ono.roberto@gmail.com ORCID [0009-0006-8650-629X](https://orcid.org/0009-0006-8650-629X)

Version 2, 5 October 2026 (arXiv:2608.28886v2; Zenodo [10.5281/zenodo.23163706](https://zenodo.org/record/23163706)), with the
corrections listed after the abstract.

Version 1: 28 August 2026 (arXiv:2608.28886v1; Zenodo [10.5281/zenodo.22714136](https://zenodo.org/record/22714136)).

Abstract

What does a generation loop gain from learning on its own verified successes? In cycles of generate, verify, select and LoRA-consolidate on online bin packing, training on value-filtered candidates shifts what the model writes on held-out variants toward value (-1.7 points of excess, $p = 0.008$; -3.1 against a random-consolidation control, $p = 0.004$) while the best observed candidate converges to the classic heuristic’s level and no further. A confirmation battery replicates the whole procedure three times, with fresh seeds and a never-consulted held-out set read exactly once: the mean was nearly identical in all three lineages (-2.0 , -1.8 , -1.9), and after aggregating within held-out variant all seven evaluable variants favored consolidation ($p = 0.008$). The best observed candidate moved *to* the classic heuristic’s level, exactly (0.021028, the lowest variant’s classic level, in all three lineages, for attract and for the random control alike); attract never went beyond its variant’s classic level, while the random control did in 7 of 24 lineage-variant cells. A matched SFT-only control shows the supervised anchor, not repulsion from bad candidates, does the concentrating (96% of candidates land exactly at the classic heuristic’s level). On far families the tails cut both ways: consolidation lowers the per-candidate rate of better-than-classic candidates (10% to 3.9%) while its larger production yields more such candidates absolutely (5 against 1, on few events). As motivation we report the inference-time ledger that led here: a model-written schematic recap buys judged document integration and nothing buys development; a verifier written into the stream is imitated, 16.4 fabricated verdict lines per notebook. In this arrangement (an 8-bit model in MLX, five cycles), mean quality among valid candidates can be bought and replicated; under attraction the observed best goes to the classic and, so far, never beyond it.

Changes in version 2. An audit of version 1 against the run data and the notebook found errors, corrected here. (i) The replication’s “no beyond-classic candidate anywhere” was wrong for the random control: 14 of its 529 valid candidates beat their variant’s classic level (7 of 24 lineage-variant cells, margins 0.0009–0.0037), as did 1 of the base’s 147; attract’s 432 did not (abstract, Sections [1](#), [2](#), [6](#), [7](#), [9](#), Table [4](#)). (ii) The SFT-only adapter’s far-family mean was the pooled value (0.1278); per family it is 0.0786, close to the anchored arm’s, so “generalizes worst” is withdrawn. (iii) Smaller statements and numbers: strict finds before the final cycle (rare, not absent); the captions of Table [2](#) and Figure [1](#); the anchored arm as one of four with no far better-than-classic candidate; per-variant and per-family values in three rows of Table [3](#); the base’s far yield per notebook (0.06); three ranges; “favoring attract” for negative differences; registration dates, with an SFT-only row (Table [5](#)). (iv) Version 1 placed all run data and judgments in the repository, but those of batteries M, M+, S and V were not public; they are now in an evidence deposit, and the judging dates are 18–20 August. (v) Added: a limitation on a later bf16/PyTorch run, a statement on AI assistance and on who read the drafts, and bibliography fixes.

1 Introduction

Ono Filho [7] dismantled a cognitively inspired generation loop over 24 conditions and found that most of its effect on judged novelty lives in one operation, the *interruption*: a new subject injected every few hundred tokens into a stream whose literal repetition is damped. On windows of fresh generated text the interruption raises judged surprise by 1.2 to 1.4 points over habituation alone, on three base models, a post-trained one and two genres, and a pre-registered replication confirms it. The same study closed with three negatives. Read as whole documents, none of the streams builds anything (no arm above 2.5 on integration, development, coherence or surprise, 0–10, judged on 4,500 tokens with the injected text removed); a judge-gated review that decides when not to interrupt does not change that; and on online bin packing, where a deterministic verifier replaces the judge, the interruption multiplies the valid, distinct candidate heuristics a base model writes three- to fourfold without raising the quality of the best. The interruption is a variation operator, and variation without selection is organized noise.

This paper takes up the question that study left: what would make local gains compose? We proceed by giving the loop, one piece at a time, the things a mind that discovers has and the loop did not, and measuring each under the protocol that the companion study showed to be necessary (generated-only windows, fresh-only estimates, a document-level reading, the premise as the unit of inference, a deterministic verifier where one exists). The pieces are (i) a *schematic* memory, in which the model carries a compressed recap of its own stream instead of the verbatim text that entrains it [1]; (ii) a *standing question*, in which the interruption points into the stream’s own unresolved matter rather than away from it [16]; (iii) *selection*, a FunSearch-style evolution over the candidates the interruption multiplies [9]; (iv) a *verifier inside the stream*, whose verdict on each candidate is written back into the notebook; and (v) *consolidation*, STaR/ReST cycles in which the model is fine-tuned (LoRA) on its own training-verifier-selected candidates and generates again [3, 11, 17], followed by preference-based repulsion [8] and by quality-diversity selection [4, 5] designed to keep the tails. The first two pieces are reported briefly as the study’s motivation; the paper’s core is the consolidation thread and its replication. Every battery was pre-registered before it ran; the pre-registrations, their dated amendments and the results are in the laboratory notebook released with the code.

Contributions. (1) Consolidation by LoRA on training-verifier-selected candidates moves the candidate distribution on never-seen variants of the trained family toward value ($p = 0.008$; against a random-consolidation control, $p = 0.004$), moves the observed best to the classic heuristic’s level and not beyond it, and does not transfer to other families. (2) A confirmation battery: three independent lineages with fresh seeds and a never-consulted held-out set read once find the same mean shift in each ($-2.0, -1.8, -1.9$ points; all seven evaluable variants favor consolidation, $p = 0.008$) and the same convergence of the best, exactly at the classic level in every lineage; a matched SFT-only control shows the supervised anchor, not the repulsion term, does the in-family concentrating, and out of family the two anchored adapters are close (0.0786 vs 0.0758). (3) The tails, both ways: on the far families the per-candidate better-than-classic rate is highest in the untrained base while attraction’s larger production yields more absolute tail events; pass@ k against the classics is zero for base, attract and random on the first held-out set at cycle five. (4) As motivation, the inference-time ledger: a model-written schematic recap gives the program’s highest document integration at no local cost, a verifier written into the stream is imitated rather than used (16.4 fabricated verdict lines per notebook), and no injection moves development past the progression wall. (5) A reading of the whole: inference-time stimulation reorganizes a fixed prior; weight consolidation moves its mass toward the known good, with the per-candidate cost in the tails and the production gain reported side by side.

2 Related work

Self-training on one’s own outputs. STaR [17] fine-tunes a model on its own rationales that led to correct answers; ReST [3] and ReST-EM [11] alternate sampling, filtering by a reward and supervised fine-tuning, and report gains that saturate after a few iterations. Our attract arm is that recipe at office scale (an 8B model, forty examples per cycle, a rank-8 adapter) with three additions: a same-size random-consolidation control, which separates “trained on its own outputs” from “trained on what the verifier approved”; transfer measured on variants and families that never entered any training set; and the rate of candidates better than the classic heuristics, which is where the cost of self-training shows.

Our consolidation result is close to the finding that reward training concentrates probability on existing trajectories, raising low- k performance while large- k coverage stays or falls [15]. The literature also contains the counterexamples that sharpen its scope: boundary-guided curriculum RL reports raising both pass@1 and pass@256 above the base [2], and representation-based exploration bonuses improve pass@ k in post-training [13]; both add what is absent here, external guidance or directed exploration beyond the model’s own rollouts.

Preference optimization. DPO [8] optimizes a policy against a frozen reference from chosen/rejected pairs. Our repulsion arms use it with the model’s own worst candidates, or the clones of the known optimum, as rejected examples; unanchored it degenerates, anchored with a supervised term on the chosen side it sharpens the policy onto the known optimum. Both behaviours are known failure and success modes of preference optimization; what we add is their effect on the tails of a candidate distribution measured by a verifier.

Quality-diversity. MAP-Elites [5] and novelty search [4, 12] keep the best solution per behavioural niche rather than the best overall. Our qd arm consolidates one elite per niche of the verifier’s per-instance behaviour profile; it keeps functional diversity and, in this regime, loses value and the tails alike.

Verified search. FunSearch [9] and AlphaEvolve [6] pair an untrained proposal model with a hard evaluator and many samples, without fine-tuning the proposer. Our results are consistent with that choice: on the far families, the untrained prior has the highest rate of better-than-classic candidates, and every consolidation we ran lowered it.

Memory, tension, and the judge. Schematic memory as reconstruction [1]; unfinished tasks as retained tension [16]; compression progress as intrinsic reward [10]. On the instrument, the companion study documents what a windowed LLM judge cannot see (injected text, self-copy beyond its horizon, windows against wholes); the fabricated verdicts of Section 5 are the same entrainment acting on a verifier’s output placed in the prompt.

3 Setup

Generators and loop. Unless stated, the generator is Qwen3-30B-A3B-Base (mixture of experts, 8-bit, MLX on an Apple M5 Pro); battery C uses Qwen3-8B-Base (8-bit) because it is fine-tuned. A cell is one premise (the ten premises of the companion study) continued for 4,500 tokens with the end-of-text token masked; habituation (windowed repetition penalty, window 512, factor 1.15) is on in every arm; interruptions fire on a clock every 300 tokens. Arms differ in what is injected and in whether the context is kept or rebuilt.

arm (Qwen3-30B-A3B-Base, period 300)	integration	development	coherence	surprise
habituation, no interruption	1.40	1.10	1.50	1.40
verbatim context (clock300)	0.80	0.40	1.50	0.50
reset + new subject	2.00	1.60	2.20	2.50
schema: reset + recap + new subject	2.70	1.80	2.30	2.70
question: preserved + open question	1.50	1.10	1.90	1.00
agenda: reset + recap + open question	2.50	2.00	1.90	2.60

Table 1: Document-level reading (whole 4,500-token stream, injected text removed; Opus 5, $k = 3$; cell means over ten premises, 0–10). The schematic recap gives the highest integration of the program; no arm passes 2.0 on development.

Judging. Windows of generated text only (96 tokens, starting 32 after any injected text, never crossing an injection), judged by Claude Opus 5 ($k = 5$, median) on surprise, connection and coherence with 600 tokens of context; windows flagged as self-copy (half their 12-token shingles seen earlier in the stream) are excluded from the fresh-only estimates. The whole stream, injected text removed, is judged once per cell ($k = 3$) on integration, development, coherence and surprise. The premise is the unit: cell means, bootstrap intervals over cells, exact sign-flip permutation tests on paired cells, one-sided where pre-registered.

The problem with a verifier. Online bin packing in the FunSearch form: the model writes a notebook of `priority(item, remaining)` functions; the verifier packs instances and reports the mean excess over the lower bound. Batteries S and V use the ten distribution variants of the companion study (item ranges within $[0.1, 0.8]$, where best fit is near-optimal); battery C moves to a family of small-item distributions in which best fit is measurably beatable (Section 6).

4 Memory and tension: the inference-time motivation

This section is the study’s motivation, kept brief; the paper’s core is the consolidation thread that follows. Five arms varied what an interrupted stream remembers and what the interruption asks (period 300, ten premises): *verbatim* keeps the full context across interruptions; *reset* rebuilds it from the premise and the new subject; *schema* rebuilds it from the premise, a two- to three-sentence recap the model writes about its own last 1,200 tokens, and the new subject; *question* keeps the context and injects the open question the model extracts from its own stream; *agenda* combines reset, recap and question. Recaps and questions count as injected text (never judged; removed from the document), with an anti-copy ladder for the model’s tendency to repeat its own recap.

Pre-registered contrasts (one-sided exact paired permutation, ten premises): schema beats verbatim on the composite by $+1.65$ [$+1.15, +2.25$], $p = 0.001$; against reset, the registered composite is not supported ($+0.45$, $p = 0.13$; decomposed, integration $+0.70$ [$+0.10, +1.30$], development $+0.20$), so most of the headline is the reset and the recap adds integration on top, at no cost on fresh windows (all fresh contrasts n.s., self-copy 0%). The question beats the neutral subject under preserved context ($+0.70$, $p = 0.004$) but matches habituation on development; the agenda adds nothing over schema ($+0.20$, $p = 0.38$; the document judge: “restarts the premise over and over with new casts and genres”). The reading that motivates the rest of the paper: remembering the schema beats re-reading the text, but every registered injection moves integration and never development above about 2/10, the *progression wall*.

5 Selection and the verifier

Battery S: the interruption inside a selection loop. A FunSearch-style evolution (10 variants $\times$ 8 generations $\times$ 16 samples per generation; the prompt shows the current population; paired by variant) with and without the interruption’s angle line injected into the prompt. S1 (gain of the champion on held-out instances): +0.0013 [-0.0015, +0.0047], $p = 0.38$. S2 (cumulative distinct valid candidates): 89.5 vs 95.1, $p = 0.86$ against the hypothesis. S3 (generation of first escape from best fit): +0.5, n.s. Under selection the interruption adds nothing, not even the diversity it added in a single stream: the population prompt already supplies what the interruption supplied.

Battery V: the verifier inside the stream. The notebook of the companion study’s battery B, generated in chunks; whenever a `priority` function closes, the harness scores it on the training instances and writes the verdict into the notebook as a comment (“# Verifier: `priority_x` scored mean excess 0.0812 on the training instances (worse than best fit, 0.0752); no improvement.”). A second arm adds, every 300 tokens, a deterministic scoreboard (“# So far: N functions tried; the best is ...; best fit sits at ...”) and a value-anchored open question. Ten variants $\times$ two seeds; baselines are battery B’s plain and angle arms re-scored in text order. V1 (held-out gain, feedback vs angle): +0.0003, $p = 0.48$. V2 (candidates improving on the notebook’s own training best, the verifier’s measure of progression): 0.1 vs 0.3 per notebook, in the opposite direction. V4 (scoreboard adds to feedback): nothing; the scoreboard arm writes fewer functions than every other arm (1.1 per notebook). Progression is near zero in every arm (0.1–0.3 improvements per notebook), with the standing caveat that best fit is near-optimal on these distributions.

Fabricated verdicts. The model wrote on average 16.4 “# Verifier:” lines per notebook in the feedback arm (6.2 with the scoreboard) against about three real ones: verdicts with invented scores and “new best of this notebook” that the harness never issued. Injecting the *format* of a reward signal into the context makes a base model continue a highly predictable textual pattern present in its context; whether the fabrication functions as reward or as mere format imitation is untested (no fabricated verdict was produced by the harness; their correlation with true scores and their effect on subsequent candidates are open analyses). The rule is narrower than “keep verdicts out of the prompt”: do not place forgeable verifier syntax in the same generative channel without provenance (a non-generable channel, a nonce, reserved tokens) and without measuring imitation.

6 Consolidation: moving the prior

Design (pre-registered 2026-08-21, extensions 2026-08-21 and 2026-08-22). Qwen3-8B-Base (8-bit), QLoRA (rank 8, 16 layers, lr 10^{-5} , batch 2, prompt masked, ≈ 4 epochs over the consolidation set). Domain: a family of small-item bin-packing distributions in which best fit is beatable by simple heuristics (headroom 0.3–0.5 points of excess; first fit ties or beats best fit), ten training variants and eight held-out variants that never enter any training set; six far families (Weibull, triangular, bimodal, OR-library) for far transfer; verification on 20×200 -item instances; a *find* beats the better of best fit and first fit by 0.001 on the variant’s training instances. Cycle 0 is a shared base generation (3 notebooks $\times$ 1,200 tokens per variant, an angle comment every 250 tokens). Cycles 1–5: *attract* trains a fresh adapter on the finds plus the top-40 candidates by training excess from its own lineage; *random* trains on a same-size random sample of valid candidates (the control that separates “trained on itself” from “trained on what is worth”); *base* generates without an adapter at the same seeds. On the cycle-5 lineage we then

arm	cycle	candidates	mean excess	rel. to min(BF,FF)	best	distinct	copies
base	0	65	0.0633	+0.0322	0.0408	0.99	–
base	1	53	0.0562	+0.0251	0.0327	1.00	–
attract	1	235	0.0559	+0.0248	0.0319	0.98	10%
random	1	276	0.0708	+0.0397	0.0332	1.00	5%
base	2	42	0.0583	+0.0272	0.0317	0.97	–
attract	2	198	0.0430	+0.0119	0.0312	0.99	6%
random	2	274	0.0697	+0.0386	0.0310	0.99	8%
base	3	55	0.0644	+0.0333	0.0414	1.00	–
attract	3	214	0.0516	+0.0204	0.0311	0.99	7%
random	3	334	0.0653	+0.0342	0.0312	0.99	1%
base	4	57	0.0532	+0.0221	0.0323	1.00	–
attract	4	195	0.0420	+0.0109	0.0311	0.98	16%
random	4	276	0.0510	+0.0199	0.0307	1.00	0%
base	5	47	0.0556	+0.0244	0.0325	1.00	–
attract	5	226	0.0386	+0.0074	0.0311	1.00	1%
random	5	287	0.0695	+0.0383	0.0311	0.99	0%

Table 2: Battery C on the eight held-out variants (never in any training set; test instances): candidates per arm and cycle, mean test excess of valid candidates (cell means over variants), the same relative to the better of best fit and first fit, the best candidate, the distinct/valid ratio and the share of literal copies of the arm’s training set. The value-filtered adapter lowers the mean; its best reaches the classic level and stays there.

trained: *repel* (DPO, $\beta = 0.1$, 80 steps, chosen = the attract set, rejected = the worst valid candidate of the same variant), *repel-anch* (the same plus a token-mean supervised term on the chosen side, $\alpha = 1$), *repel-early* (plain DPO, 20 steps), *qd* (supervised, one elite per behavioural niche of the verifier’s per-instance profile, $k = 3$ niches per variant plus one seat for the best clone of the classic attractor) and *repel-mode* (anchored DPO whose rejected examples are the best-fit clones themselves). Hypotheses, on the held-out variants: C1 (primary), attract < base on the mean test excess of valid candidates (the prior moved on variants it never saw); C2, attract < random; C3, higher find rate; C4 (two-sided), the diversity cost; R1’ repulsion adds to attraction; R2’ the best moves; T1 the tail rate (share of candidates better than the classics) on far families; T2 functional collapse.

Attraction. Table 2 gives the held-out reading by arm and cycle (Figure 1). At the pre-registered final cycle (3), C1 was -0.0129 $[-0.0317, +0.0027]$, $p = 0.109$, not supported at $\alpha = 0.05$, with the direction consistent (7 of 8 variants better) and the cycle-2 contrast supported (-0.0153 , $p = 0.016$); a pre-registered extension to five cycles closed it: -0.0170 $[-0.0268, -0.0079]$, $p = 0.008$, the attract arm ending at 3.86% mean excess, 0.74 points above the better classic heuristic, against 5.3–6.4% for the base. C2: attract vs random, -0.0138 ($p = 0.035$) at cycle 3 and -0.0309 ($p = 0.004$) at cycle 5: the value filter matters. The random control hurts at cycle 1 (+1.5 points over base: consolidating mediocre functions makes the next ones worse) and is neutral to harmful later. C3: strict finds are rare (at most 1.7%, attract at cycle 4) and zero in every arm at the final cycle, and the best candidate per variant is nearly the same in every arm ($\approx 3.1\%$ excess, the level of best fit): the ceiling did not move. C4: no collapse over five cycles; the distinct/valid ratio is 0.97–1.00 throughout, literal copies of the training set 1–16% (attract); both adapters make the model write 3.4–6.5 $\times$ as many valid functions per notebook as the base. Far transfer: on six distributions of other families the final attract adapter equals the base on the mean ($+0.001$, $p = 0.69$; $n = 5$ of the six far families; the base arm produced no valid candidate in one) and beats the random adapter by -0.020 ($p = 0.016$); the random adapter is worse than the base there too ($+0.017$). What moved is the prior for the




Figure 1: The title of this paper, as data (held-out variants, never in any training set; cycle 0 is the shared pre-consolidation base). Left: mean excess over the lower bound falls under attraction and nowhere else. Right: from the second cycle the trained arms’ best stays within 0.0005 of the 0.0311 classic level, while the base’s best fluctuates above it (0.0317–0.0414). Consolidation moves the mean; attraction reaches but does not exceed the classic level, while the random control dips below it at cycles 2 and 4.

adapter (cycle-5 lineage)	held-out variants (8)				far families (6)	
	cand.	mean exc.	at classic	levels	mean exc.	> classics
base (none)	47	0.0556	20%	0.81	0.0838	10.0%
random	287	0.0695	16%	–	0.1004	0.0%
repel-early (DPO, 20 steps)	71	0.0721	34%	–	0.0994	2.4%
attract (SFT on value)	226	0.0386	72%	0.14	0.0803	3.9%
repel-anch (DPO + anchor)	144	0.0311	100%	0.07	0.0758	0.0%
sft-only (anchor term alone)	78	0.0319	97%	–	0.0786	0.0%
qd (SFT on niche elites)	293	0.0598	25%	0.33	0.1012	0.6%
repel-mode (DPO, from attractor)	62	0.0792	2%	0.50	0.0935	0.0%
repel (DPO, unanchored)	13	0.1056	0%	–	–	–

Table 3: The ladder of adapters. Held-out: mean test excess of valid candidates, share scoring exactly at the level of the better classic heuristic, distinct score levels per valid candidate; far families: mean excess and share of candidates better than the classics. Means and shares are averaged over variants (held-out) or families (far), as in Table 2; candidate counts are totals. The mass walks onto the known attractor; attraction lowers the far-family tail rate per valid candidate while raising absolute tail yield through production. Rows below attract are single-lineage mechanistic follow-ups.

family it was trained on. The adapter of the last cycle was taught mostly tight-bin functions (24 of 40 of the form $1/(r**2+eps)$, 13 best-fit-like): consolidation pulls the mass toward the best attractor the verifier has approved.

Repulsion. Unanchored DPO (loss 0.69 $\rightarrow$ 0.0015) degenerated: 1.1 closed functions per notebook (attract 9, base 2), 60% comment lines and recursive non-code; the few valid candidates were far worse than the base (10.6% vs 5.6% excess). Pushing the policy away from bad functions, unanchored, pushes it away from writing functions. Anchored (the same loss plus a token-mean supervised term on the chosen side; pre-registered), the repulsion writes again (144 valid candidates, +12 per cell over the base, $p = 0.02$) and adds to attraction: mean held-out excess 3.11% against 3.86% (-0.0074 $[-0.0104, -0.0047]$, $p = 0.008$), which is exactly the level of the better classic heuristic. The best candidate does not move (0.0311 in both arms), and the reason is the finding: **every** held-out candidate of the anchored-repulsion adapter scores exactly at the best-fit level (all are monotone “fullest bin” variants), against 72% for attraction,




Figure 2: What consolidation costs. Left: the share of candidates *better* than the stronger classic heuristic on the far families, the tail where discovery lives, shrinks from 10% (base) to 3.9% (attraction) to at most 0.6% under every tail-preserving objective we built. Right: on held-out variants the consolidated arms concentrate *exactly* at the classic’s level (attraction 72%, anchored repulsion 100%, a functional collapse), against 20% for the base. The bar heights are per-family average rates (the registered estimand); the counts above them are pooled events over all candidates, a different aggregation (base: 10.0% macro but 1/18 = 5.6% pooled; attraction: 3.9% macro, 5/156 = 3.2% pooled). Both rest on 0–5 events per arm and are preliminary. Per notebook, attraction’s absolute tail yield is higher than the base’s (0.28 vs 0.06). Competence per candidate is bought with idiosyncrasy; total yield depends on production.

16–20% for the base and the random control, 34% for an early-stopped plain DPO (otherwise worse than attraction, +0.034, $p = 0.008$) and 0% for the degenerate one (Table 3). Hash-level diversity survives (1.00); functional diversity does not. On the far families the anchored adapter is the best arm on the mean (7.6% vs 8.4% base) because “be best fit” is good everywhere, and, with the random control, the SFT-only anchor and mode repulsion, one of four arms with no candidate better than the classics (Table 3).

Equal-budget comparisons. At matched candidate counts the picture sharpens rather than reverses. Best-of- k over 2,000 subsamples per variant: at $k=10$ valid candidates, attraction’s subsample best (0.0311) beats the random control’s (0.0381) in eight of eight variants each, while the base reaches ten valid candidates in only one variant of eight (0.0449 there); at $k=40$ only the random arm has coverage at all. The base’s small-sample best advantage is therefore in part a production deficit, and no arm’s pass@ k against the classics ever leaves zero. The pooled quantiles say the same about the whole distribution: attraction dominates the base at p10, p50 and p90 (0.0257/0.0283/0.0676 against 0.0285/0.0483/0.1075), so the shift is of the distribution, not only of its mean.

A two-part reading. Because the headline measure conditions on validity, the effect decomposes into two parts that the replication battery lets us separate: consolidation multiplies the production of valid candidates (base 33–63 per lineage against attraction’s 108–184, with one base variant producing none at all) and, conditional on validity, shifts their mean quality (−0.019). “Moving the mean” is the second part; the first may matter as much in practice, and both replicate.

The registered SFT-only control (a single-lineage mechanistic follow-up, like the repulsion, QD and far-family results that follow). The anchored objective combines repulsion from the worst candidates with a supervised term on the chosen ones; which half concentrates the policy? The matched control (same pairs, same 80 steps, same learning rate and data order, the DPO term removed) answers: the supervised anchor alone puts 75 of its

lineage	arm	candidates	valid	mean excess	best	Δ vs base
1	base	71	63	0.0612	0.0247	
1	attract	201	184	0.0416	0.0210	-0.0196
1	random	219	191	0.0619	0.0210	+0.0007
2	base	79	51	0.0605	0.0243	
2	attract	162	140	0.0430	0.0210	-0.0176
2	random	232	205	0.0548	0.0210	-0.0057
3	base	41	33	0.0613	0.0243	
3	attract	120	108	0.0424	0.0210	-0.0188
3	random	202	133	0.0549	0.0210	-0.0064

Table 4: The replication battery on the never-consulted held-out set (means averaged over the seven common variants; Δ = per-lineage variant-averaged difference to base; best = minimum over all held-out variants). Three independent runs of the whole procedure produce three nearly identical effects on the mean and the same movement of the best: to the classic level, in every arm that trains, and never beyond it under attract.

78 held-out candidates (96%; 97% as a mean over variants, Table 3) exactly at the classic level (anchored: 100%; attraction: 72%), at mean excess 0.0319 against the anchored arm’s 0.0311. Repelling the bad was never the mover; supervised attraction to the chosen does nearly all the concentrating, and the repulsion term adds the last few points. On the far families the SFT-only adapter averages 0.0786 (per-family mean; 0.1278 pooled, dominated by one family), close to the anchored arm’s 0.0758 and below the base’s 0.0838, with no better-than-classic candidate: the supervised anchor drives the concentration in and out of the family; the DPO term adds little out of it.

Independent replication across three lineages. The confirmation battery re-ran the whole procedure three times: independent lineages with fresh generation and selection seeds, five cycles fixed in advance, and a second held-out set of eight variants, disjoint from every earlier set, generated and read exactly once after cycle five. Table 4 and Figure 3 give the design’s answer. The mean moves in every lineage by nearly the same amount (-0.0196 , -0.0176 , -0.0188 ; pooled -0.0187 [-0.0308 , -0.0080], aggregated by held-out variant, favoring attract in all seven evaluable variants, exact sign-flip $p = 0.008$, hierarchical bootstrap over lineages and variants); attract beats the random control in each (-0.0149 pooled, $p = 0.016$); and the best converges in all: the adapters’ best lands exactly at the classic level (0.021028, that of the lowest variant) in each lineage. Per variant, attract never passes its variant’s classic level, while the random control does with 14 candidates (7 of 24 lineage-variant cells, margins 0.0009–0.0037) and the base with one. Production replicates too: the base wrote 33–63 valid candidates per lineage against the adapters’ 108–205, and in one variant of one lineage the base produced no valid candidate at all (the eighth variant is therefore excluded from the paired contrast, a conditioning we flag rather than hide). The registered shared-pool one-shot (top-40 vs random-40 from a single base pool) points the same way (-0.008 , $p = 0.078$ at $n=8$, underpowered alone). Two readings sharpen the picture. First, a sensitivity for the excluded variant: restoring it under imputation (assigning the base its best or its worst observed variant mean there) gives $n=8$ with all eight variants favoring attract, $p = 0.0039$ under either bound, so the conditioning was conservative. Second, a three-effect decomposition: the random control also reaches the best of 0.021028 in every lineage, so generic self-training plus increased valid production suffices to *reach* the classic level (and, in seven lineage-variant cells, to pass it by at most 0.004); value filtering is what the distributional mean shift requires.

Keeping the tails. Two objectives built to keep the tails (pre-registered) recovered functional diversity and nothing else (Figure 2). The qd adapter has 25% of candidates at the classic level




Figure 3: Independent replication, variant by variant: each line pairs one held-out variant’s mean under base and under attract, one panel per lineage; the right panel shows the best candidate. The slope repeats across lineages; the best moves to the classic level in all three lineages and stops there.

(attract 72%, $p = 0.004$) and $2.4\times$ the distinct score levels, at a cost of +2.1 points of mean excess ($p = 0.008$), and its better-than-classic rate on the far families is *lower* than attraction’s (0.6% vs 3.9%; T1, primary, not supported). Repulsion from the attractor itself removes the collapse entirely (2% at the classic level) and with it most of the value (7.9% mean excess, best 4.8%, worse than the base; no far tail). The diversity these objectives keep is diversity of the mediocre. Across every consolidation we ran, the per-valid-candidate rate of better-than-classic candidates on new families is highest in the untrained base (10% as a mean over families; pooled, 1 of 18 candidates) and falls with training (attract 3.9%, pooled 5 of 156; every other adapter $\leq 2.4\%$, pooled 0–1 events). The absolute yield points the other way: because consolidation multiplies valid production, attraction delivered five better-than-classic candidates against the base’s one, 0.28 against 0.06 per notebook, across two families against one. Both readings rest on zero to five events per arm and are preliminary; what they jointly say is that consolidation concentrates each candidate toward the known good while producing many more candidates, so the per-candidate tail rate falls even where the total tail yield does not. The anchored objectives, by contrast, produced no better-than-classic candidate at all.

7 Discussion

What can be bought. Across the two studies the ledger is now complete for the interventions we could build. Variation is cheap: the interruption raises fresh surprise by more than a point and multiplies valid candidates three- to fourfold. The document judge assigned higher integration to a schematic recap carried across resets than to any other arm, although the pre-registered composite against reset alone was not supported. Mass can be bought with the right training: a few cycles of consolidation on what the verifier approved move the candidate distribution of a never-seen variant onto the known good, and an anchored repulsion from the bad moves it there exactly. Two things nothing we tried produced: measurable document progression (three injections, selection, a verifier in the stream and a judge-gated review all leave development near 2/10) and, under the evaluated bin-packing budgets, a held-out candidate measurably beyond the classic from any value-filtered or anchored adapter (five cycles of attraction, four kinds of repulsion and a quality-diversity consolidation leave the best observed held-out level at the classic or above it, none passing it by more than 10^{-4} ; the held-out candidates beyond it, by 0.0009–0.0037, came from the untrained base and the random control).

The bit distance (a conceptual model consistent with, not demonstrated by, these data). A generative model is a distribution, not a store; the question is not whether an answer is “inside” it but how much probability mass sits near it, how many bits of search separate the

prior from the target. Inference-time stimulation reorganizes the surface of a fixed prior: it buys variation and, with a schematic memory, integration, but it does not move the mass. Weight consolidation moves the mass, and battery C shows its present reach: toward the best thing the verifier had already approved, within the family it was trained on, and at the cost of the tails where the rare better-than-classic candidates appear; objectives built to keep the tails kept diversity of the mediocre instead. This is also why the known successes of verified search [6, 9] pair an untrained proposal distribution with a hard evaluator and many samples rather than fine-tune the proposer. For discovery our results argue for a portfolio, the broad untrained prior to propose and the consolidated one to exploit, and for spending on selection.

Design rules. For long open-ended generation: damp repetition; interrupt with something new each time or reset the context; carry a compressed recap rather than the text; do not expect the whole to advance from any of this. For search with a verifier: spend on sampling and selection, not on stimulation; keep the verifier’s verdicts out of the prompt or measure the fabrication they induce; if reliability is the goal, consolidate on what the verifier approved (and never without the filter); whether a portfolio that keeps the untrained prior in the loop beats pure consolidation for discovery depends on the cost function (per candidate, the base finds more; per notebook, attraction does) and awaits a fixed-budget comparison.

8 Limitations

Ten premises and ten variants per battery; 8–30B base models at 8 bits; the document judge is a single LLM family at $k = 3$ and its absolute levels are low for every arm, so “integration 2.7” is a relative statement; the bin-packing distributions of batteries S and V have little headroom, which bounds what selection and feedback could show there (battery C’s family was chosen for headroom for that reason, and its headroom is still small: 0.3–0.5 points); the oracle that writes recaps and questions is the generator itself and sometimes repeats itself, which the anti-copy ladder mitigates but does not remove; the original battery C is one training lineage of five cycles of forty examples on an 8B model with a rank-8 adapter, whose eight held-out variants replicate the adapter, not the procedure; the replication battery (three independent lineages, new generation and selection seeds, five cycles fixed in advance, a second held-out set generated and read once) closes that gap for the mean and the value filter, on seven of eight registered variants (the base produced no valid candidate in one, the production deficit again); the original lineage’s primary contrast missed α at the registered third cycle ($p = 0.11$) and closed only in a sequential-analysis sense at the extension to five ($p = 0.008$), the held-out set having been consulted at both looks; its base arm varies by seed; candidate counts differ by arm (42–334), so the observed best is not an equal-budget comparison of tails (the matched- k analyses in the results address this); the far families are six (paired $n = 5$) and the tail rates rest on tens of candidates; the QD arm is a small implementation (three niches per variant), so “tail-preserving objectives keep the mediocre” is a statement about this implementation, not the class [14]; after the first cycle attract and random are two whole adaptive procedures rather than two selection criteria applied to one pool (the registered shared-pool one-shot control, top-40 vs random-40 from a single base pool, points the registered way, -0.008 , $p = 0.078$ at $n=8$, underpowered alone); excluding the far family in which the base produced no valid candidate treats a production failure as missing data, so the paired far contrast is conditioned on validity; the selection criterion admits lucky candidates (training finds lose on test), which is why the primary measure is the mean and not the best. A later run (laboratory notebook, 2–3 September 2026; the appendices APPENDIX_T2.md and APPENDIX_T2_8B.md are in the evidence deposit) repeated the consolidation on the same domain in bf16 on PyTorch with three cycles (Qwen3-14B-Base and Qwen3-8B-Base) and reproduced neither the base’s far-family tail rate

(0–5% per cycle, against 10% here) nor attraction’s mean shift on the held-out variants (attract minus base between -0.9 and $+0.6$ points per cycle); the claim that mean quality can be bought and replicated holds for the 8-bit MLX arrangement with five cycles, and we have not separated which of backend, precision, model size or number of cycles it depends on. The negative results are resolution-bounded nulls, not proofs of absence; the tail finding in particular deserves replication at a larger sample and on a domain with more headroom.

9 Conclusion

Given memory, tension, selection and a verifier, the interrupted loop composes better and advances no further. A schematic recap raised judged integration relative to reset, while the pre-registered composite contrast was not supported; a standing question outperformed the fixed subject change under preserved context but did not improve development over habituation alone or over schema; selection and in-stream verification showed no detectable benefit on the near-saturated variants tested, and the latter teaches the model to imitate verdicts. Consolidating training-verifier-selected candidates into the weights is the one intervention that moved what the model writes on problems it had not seen, and it moved it toward what it already knew to be good; the combined anchored objective moved it all the way there and no further, and the registered SFT-only control showed the supervised half does nearly all of that moving (96% at the classic level without any repulsion term); keeping the tails kept the mediocre. In this lineage the untrained base had the highest per-candidate better-than-classic rate on the evaluated far families, while attraction’s larger production delivered more absolute tail events; the best statistic improved to the classic level, and attraction produced no candidate measurably beyond the classic on the held-out variants. In this arrangement (an 8-bit model in MLX, five cycles), mean candidate quality among valid candidates can be bought; under attraction the observed best moves to the classic heuristic and not beyond it, in every lineage. And the central remaining question, whether the procedure replicates or the effect was a peculiarity of one training path, is now answered: three independent lineages, fresh seeds, five cycles fixed in advance, a second held-out set read exactly once, and the same result three times, -2.0 , -1.8 and -1.9 points of mean excess against the base (aggregated by held-out variant, favoring attract in all seven evaluable variants, $p = 0.008$), and the same convergence of the best: in every lineage the adapters’ best candidate lands exactly on the classic heuristic’s level (0.021028), and attract produces no beyond-classic candidate (the random control 14, the base one). The best statistic itself improves from the base’s 0.0243–0.0247 to that level, an attractor reading rather than an immobility one; the exact sign-flip on the variant-paired best difference does not resolve it (-0.004 , $p = 0.094$; the hierarchical bootstrap interval $[-0.011, -0.000]$ reflects a different aggregation and should be read as descriptive).

Acknowledgements and tooling. The experiments, analyses and text of this paper were produced by the author working with an AI assistant (Claude, Anthropic; model Fable 5 for version 1, August 2026, and Opus 5.5 for the corrections of version 2, October 2026) inside the Claude Code environment. The assistant wrote and ran the code, operated the data collection under the author’s direction and drafted the manuscript; every design decision, result and claim was reviewed by the author, who is responsible for the content. A colleague read the drafts of version 1 in five rounds between 24 and 28 August 2026, and the responses to each round are dated in the laboratory notebook; this was a colleague’s reading, not peer review. The corrections of version 2 come from a read-only audit of version 1 against the run data and the notebook, carried out by instances of the same assistant and approved by the author. The LLM judge (Claude Opus 5 on Amazon Bedrock) is part of the instrument; generators and fine-tuning ran locally on Apple silicon via MLX.

A Reproducibility and experimental provenance

battery	registered	data seen	primary	status
S (selection)	2026-08-19	paper-1 batteries	held-out gain	null
M/M+ (memory, tension)	2026-08-20	paper-1, S	schema vs verbatim	supported
V (verifier)	2026-08-20	M, S	held-out gain	null; fabr. verdicts
C cycles 1–3	2026-08-21	none	attract vs base (c3)	not supported ($p = 0.11$)
C cycles 4–5	2026-08-21	c3 result	attract vs base (c5)	sequential, $p = 0.008$
C-repel, N (tails)	2026-08-21/22	C	keep the tails	degenerate / not kept
Matched- k analyses	2026-08-24	all C	best-of- k , pass@ k	reported
C-rep (3 lineages)	2026-08-26	all above	replication	mean yes; best to classic
SFT-only control	2026-08-27	all above	at-classic share	the anchor concentrates

Table 5: Order of registrations and looks (dates from the commit history; M+ was registered after the M result). The cycle-5 result is a prospective sequential extension, not an independent confirmation; the independent lineages with never-consulted held-outs are the confirmation battery.

Code, pre-registrations and the dated laboratory notebook are in the repository of the companion study (<https://github.com/RobertoOno/interrupting-the-loop>): the oracle arms (`schema_reseed`, `anomaly_reseed`, `agenda_reseed` in `scripts/dream_run.py`), battery S (`scripts/evolve_interrupt.py`), battery V (`scripts/problem_loop.py`), battery C and its extensions (`scripts/consolidate.py`, `dpo_lora.py`, `run_c.sh`, `run_c_ext.sh`, `run_c_repel*.sh`, `run_n.sh`) and the analyses that generate the tables (`docs/APPENDIX_GEN.md`, `APPENDIX_S.md`, `APPENDIX_V.md`, `APPENDIX_C*.md`, `APPENDIX_N.md`). The candidate records of batteries C and C-rep are in the release `run-data-v1` of that repository (Zenodo [10.5281/zenodo.22714138](https://zenodo.org/record/105281)); the run data and judgments of every battery reported here, including M, M+, S and V, are in the evidence deposit [10.5281/zenodo.23163299](https://zenodo.org/record/105281). Judging used `anthropic.claude-opus-5` on Amazon Bedrock (18–20 August 2026) under the prompts of the companion study; fine-tuning used `mlx_lm` (QLoRA on the 8-bit model, peak 14 GB, about 8 minutes per adapter) and a small DPO trainer in MLX released with the code. Every hypothesis, amendment and result is dated in `docs/PLANO.md`.

References

- [1] Frederic C. Bartlett. *Remembering: A Study in Experimental and Social Psychology*. Cambridge University Press, 1932.
- [2] Pengxiang Cai, Tianchen Fang, Xiaohan Li, Qingyuan Zeng, Guocong Li, and Jintai Chen. Curriculum reinforcement learning can incentivize reasoning capacity in LLMs beyond the base model, 2026. arXiv:2606.22317.
- [3] Caglar Gulcehre, Tom Le Paine, Srivatsan Srinivasan, Ksenia Konyushkova, Lotte Weerts, Abhishek Sharma, Aditya Siddhant, Alex Ahern, Miaosen Wang, Chenjie Gu, Wolfgang Macherey, Arnaud Doucet, Orhan Firat, and Nando de Freitas. Reinforced self-training (ReST) for language modeling, 2023. arXiv:2308.08998.
- [4] Joel Lehman and Kenneth O. Stanley. Abandoning objectives: Evolution through the search for novelty alone. *Evolutionary Computation*, 19(2):189–223, 2011.
- [5] Jean-Baptiste Mouret and Jeff Clune. Illuminating search spaces by mapping elites, 2015. arXiv:1504.04909.
- [6] Alexander Novikov et al. AlphaEvolve: A coding agent for scientific and algorithmic discovery, 2025. Google DeepMind; arXiv:2506.13131.

- [7] Roberto I. Ono Filho. Interrupting the loop: Periodic subject changes raise judged surprise and connection in base language models. *arXiv:2608.19893 [cs.CL]*, 2026.
- [8] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, 2023.
- [9] Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M. Pawan Kumar, Emilien Dupont, Francisco J. R. Ruiz, Jordan S. Ellenberg, Pengming Wang, Omar Fawzi, Pushmeet Kohli, and Alhussein Fawzi. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024. doi: 10.1038/s41586-023-06924-6.
- [10] Jürgen Schmidhuber. Formal theory of creativity, fun, and intrinsic motivation (1990–2010). *IEEE Transactions on Autonomous Mental Development*, 2(3):230–247, 2010.
- [11] Avi Singh, John D. Co-Reyes, Rishabh Agarwal, Ankesh Anand, Piyush Patil, et al. Beyond human data: Scaling self-training for problem-solving with language models, 2023. *arXiv:2312.06585*.
- [12] Kenneth O. Stanley and Joel Lehman. *Why Greatness Cannot Be Planned: The Myth of the Objective*. Springer, 2015.
- [13] Jens Tuyls, Dylan J. Foster, Akshay Krishnamurthy, and Jordan T. Ash. Representation-based exploration for language models: From test-time to post-training, 2025. *arXiv:2510.11686*.
- [14] Chunhui Wan et al. LoongFlow: Directed evolutionary search via a cognitive plan-execute-summarize paradigm, 2025. *arXiv:2512.24077*.
- [15] Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang. Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model? In *Advances in Neural Information Processing Systems*, 2025. *arXiv:2504.13837*.
- [16] Bluma Zeigarnik. Das Behalten erledigter und unerledigter Handlungen. *Psychologische Forschung*, 9:1–85, 1927.
- [17] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. STaR: Bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems*, 2022.